

# Tuning Machine Learning Algorithms with mlr3

## mlr3tuning

Department of Statistics – LMU Munich

September 25, 2019



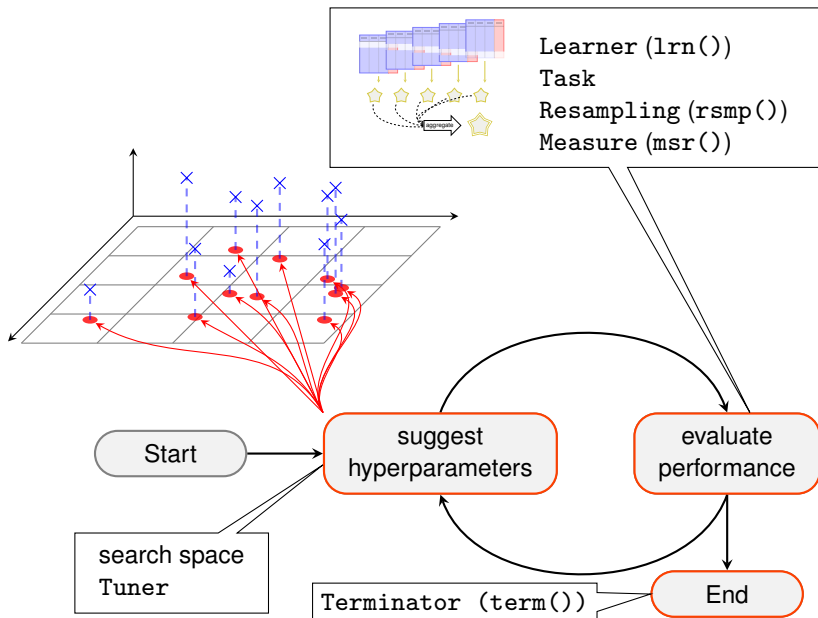
# Intro

# TUNING

- Behaviour of most methods depends on *hyperparameters*
  - We want to choose them so our algorithm performs well
  - Good hyperparameters are data-dependent
- ⇒ We do *black box optimization* (“Try stuff and see what works”)

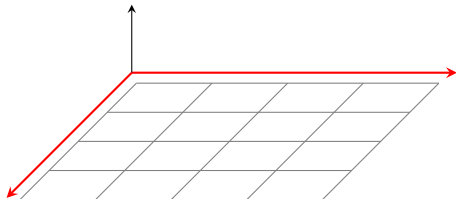
# Tuning

# TUNING



# Search Space

# SEARCH SPACE



```
ParamSet$new(list(param1, param2, ...))
```

*Numerical* parameter      ParamDbl\$new(id, lower, upper)

*Integer* parameter        ParamInt\$new(id, lower, upper)

*Discrete* parameter      ParamFct\$new(id, levels)

*Logical* parameter        ParamLgl\$new(id)


*Untyped* parameter        ParamUty\$new(id)

```
library("paradox")
searchspace_knn = ParamSet$new(list(
  ParamInt$new("k", 1, 20)
))
```

# Tuning with mlr3tuning



# OBJECTS IN TUNING



```
library("mlr3tuning")

inst = TuningInstance$new(
  tsk("iris"), lrn("classif.kknn", kernel="rectangular"),
  rsmp("cv"), msr("classif.ce"),
  searchspace_knn, term("evals", n_evals = 20)
)
```

# TUNING METHOD

- need to choose a *tuning method*
- Tuner class, load with `tnr()`, set parameters
- `gsearch = tnr("grid_search", resolution = 20)`

```
print(gsearch)

#> <TunerGridSearch>
#> * Parameters: resolution=20, batch_size=1
#> * Packages: -
#> * Properties: dependencies
```

- common parameter `batch_size` for parallelization

# CALLING THE TUNER

```
gsearch$tune(inst)
```

```
inst$result
```

```
#> $tune_x
```

```
#> $tune_x$k
```

```
#> [1] 13
```

```
#>
```

```
#>
```

```
#> $params
```

```
#> $params$kernel
```

```
#> [1] "rectangular"
```

```
#>
```

```
#> $params$k
```

```
#> [1] 13
```

```
#>
```

```
#>
```


```
#> $perf
```

```
#> classif.ce
```

```
#> 0.033
```

# TUNING RESULTS

```
ggplot(inst$archive(unnest = "params"),  
  aes(x = k, y = classif.ce)) + geom_line() + geom_point()
```



# RECAP

```
inst = TuningInstance$new(  
  tsk("iris"), lrn("classif.kknn", kernel="rectangular"),  
  rsmp("cv"), msr("classif.ce"),  
  searchspace_knn, term("evals", n_evals = 20)  
)  
  
gsearch = tnr("grid_search", resolution = 20)  
  
gsearch$tune(inst)
```

# Parameter Transformation

# PARAMETER TRANSFORMATION

- Sometimes we do not want to sample evenly from a range
- $k = 1$  vs.  $k = 2$  probably more interesting than  $k = 101$  vs.  $k = 102$

⇒ Transformations

- Part of ParamSet


Example:

- 1 sample from  $\log(1) \dots \log(100)$  (`k_before_trafo`)
- 2 transform by `exp()` in `trafo` function
- 3 don't forget to `round` ( $k$  must be integer)

```
searchspace_knn_trafo = ParamSet$new(list(  
  ParamDbl$new("k_before_trafo", log(1), log(100))  
)  
)  
searchspace_knn_trafo$trafo = function(x, param_set) {  
  return(list(k = round(exp(x$k_before_trafo))))  
}
```

# PARAMETER TRANSFORMATION

What is our transformation doing?





# PARAMETER TRANSFORMATION


Tuning again...

```
inst$result

#> $tune_x
#> $tune_x$k_before_trafo
#> [1] 2.7
#>
#>
#> $params
#> $params$kernel
#> [1] "rectangular"
#>
#> $params$k
#> [1] 14
#>
#>
#> $perf
#> classif.ce
#>          0.033
```

# PARAMETER TRANSFORMATION

```
ggplot(inst$archive(unnest = "params"),  
  aes(x = k, y = classif.ce)) + geom_line() + geom_point()
```



# Nested Resampling

# NESTED RESAMPLING

- Need to perform nested resampling to estimate tuned learner performance

⇒ Treat tuning as if it were a Learner!

- Training:

- ❶ Tune model using (inner) resampling
- ❷ Train final model with best parameters on all (i.e. outer resampling) data

- Predicting: Just use final model

- **AutoTuner**

```
optlrn = AutoTuner$new(lrn("classif.kknn", kernel="rectangular"),  
  rsmp("cv"), msr("classif.ce"), searchspace_knn,  
  term("none"), tnr("grid_search", resolution = 20))
```


# NESTED RESAMPLING

```
optlrn$train(tsk("iris"))
```

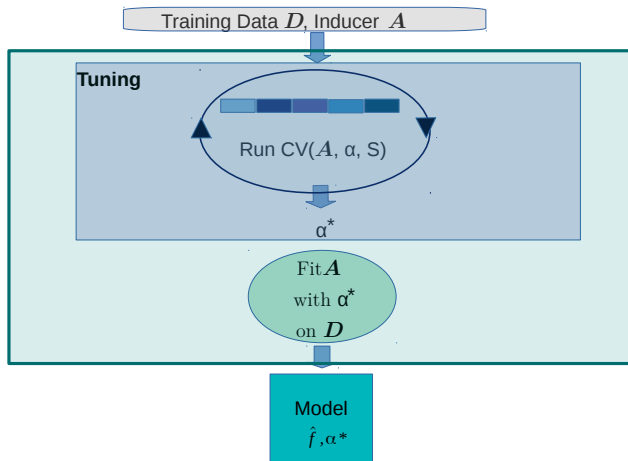
```
optlrn$model$learner
```

```
#> <LearnerClassifKNN:classif.knn>
#> * Model: data.table
#> * Parameters: kernel=rectangular, k=13
#> * Packages: withr, knn
#> * Predict Type: response
#> * Feature types: logical, integer, numeric, factor, ordered
#> * Properties: multiclass, twoclass
```

```
ggplot(optlrn$model$tuning_instance$archive(unnest = "params"),
  aes(x = k, y = classif.ce)) + geom_line() + geom_point()
```



# NESTED RESAMPLING



# NESTED RESAMPLING

```
resample(tsk("iris"), optlrn, rsmp("cv"))
```

```
#> <ResampleResult> of 10 iterations
```

```
#> * Task: iris
```

```
#> * Learner: classif.kknn.tuned
```

```
#> * Performance: 0.067 [classif.ce]
```


```
#> * Warnings: 0 in 0 iterations
```

```
#> * Errors: 0 in 0 iterations
```

# NESTED RESAMPLING

```
result = resample(tsk("iris"), optlrn, rsmp("cv"),  
  store_model = TRUE)
```

```
ggplot(result$learners[[1]]$  
  model$tuning_instance$archive(unnest = "params"),  
  aes(x = k, y = classif.ce)) + geom_line() + geom_point()
```





# Outro

# TUNING WITH MLR3TUNING

## Tuning a Learner

- 1 Construct a `TuningInstance`
  - `Task`—the Data to tune over
  - `Learner`—the algorithm to tune
  - `Resampling`—the resampling method to use
  - `Measure`—how to evaluate performance
  - `ParamSet`—the search space, possibly with `trafo`
  - `Terminator`—when to quit
- 2 Create a Tuner
  - Usually using `tnr()`
  - May have some parameters, e.g. `batch_size`
- 3 Call `tuner$tune()`

## Nested Resampling

- 1 Construct an `AutoTuner`
  - Constructor takes all arguments of a `TuningInstance` *except* `Task`
  - Also takes the Tuner as an argument
- 2 Use like a normal `Learner` in `resample()` and `benchmark()`